

Linear Models for Regression

Stylianos Sygletos


Aston Institute of Photonic Technologies - Aston University

June 10, 2021

The regression task

- It is a supervised learning task, with a goal to:
 - to predict the value of one continuous target *target* variables t
 - given a D -dimensional vector \mathbf{x} of *input* variables.
 - and based on a dataset of known inputs and outputs (training set), i.e. $\{\mathbf{x}_n, t_n\}, n = 1, \dots, N$
- From a probabilistic perspective, we aim to model the predictive distribution $p(t|\mathbf{x})$.
 - it expresses our uncertainty about the value of t for each value of \mathbf{x}
 - we can make predictions of t , for any new value of \mathbf{x}
 - we are not trying to model the distribution of \mathbf{x}
- We do not expect the predictor to be a linear function of \mathbf{x} .

An Example



Polynomial Curve Fitting with a Scalar

- With a *single* input variable

$$y(x, \mathbf{x}) = w_0 + w_1x + w_2x^2 + \cdots + w_Mx^M = \sum_{j=0}^M w_jx^j$$





where M the polynomial order, x^j the j^{th} order of x , and $\mathbf{w} = (w_0, \dots, w_M)^T$

- *What is our task?* To learn \mathbf{w} from training data $\{(x_i, t_i)\}_{i=1}^N$. We calculate the misfit between predictions $y(x_n, \mathbf{w})$ for each data point x_n and corresponding target values t_n :

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ y(x_n, \mathbf{w}) - t_n \right\}^2$$

it is zero when $y(x, \mathbf{w})$ passes through training data points.

Results with Polynomial Basis



Regression with multiple inputs

- Generalization :

- predict value of continuous target variable t given value of D input variables $\mathbf{x} = [x_1, \dots, x_D]$
- t can also be a set of variables (multiple regression)
- Linear functions of adjustable parameters
 - specifically linear combinations of *nonlinear* functions of input variables

- Polynomial curve fitting is good only for :

- single input scalar variable x
- it cannot be easily generalized to several variables


Simplest Linear Model with D inputs

- It involves a linear combination of D -input variables:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1x_1 + \cdots + w_Dx_D$$

where $\mathbf{x} = (x_1, \dots, x_D)^T$

- Known as *linear regression* because it is a linear function of the parameters w_0, \dots, w_D .
- Being a linear function of input variables imposes severe limitations
- We can extend the class of models by considering fixed nonlinear functions of input variables



Basis Functions

- If the original variables comprise the vector \mathbf{x} , then the features can be expressed in terms of basis functions $\phi_j(\mathbf{x})$:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

where $\mathbf{w} = (w_0, \dots, w_{M-1})^T$ and $\boldsymbol{\phi} = (\phi_0, \dots, \phi_{M-1})$ with $\phi_0(\mathbf{x}) = 1$ defining a *dummy basis function* for the bias w_0

- By using nonlinear basis function we allow $y(\mathbf{x}, \mathbf{w})$ to be a *nonlinear function* of the input vector \mathbf{x}

Choice of Basis Function

- ① Polynomial
- ② Gaussian
- ③ Sigmoidal
- ④ Fourier
- ⑤ Wavelets


Polynomial Basis for one variable

- Linear Basis Function Model:

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(x) = \sum_{j=0}^{M-1} w_j x^j$$

- Polynomial Basis (for single variable x)

$$\phi_j(x) = x^j$$



Disadvantages of polynomial basis :

- Global: changes in one region of input space affects others
- Difficult to formulate : number of polynomials increases exponentially with M
- Can divide input space into regions :
 - use different polynomial in each region
 - use local basis functions (e.g. *spline functions*)

Gaussian Radial Basis Functions


- Gaussian:

$$\phi_j(x) = \exp \left[-\frac{(x - \mu_j)^2}{2\sigma^2} \right]$$

- does not have probabilistic interpretation
- normalization term not needed

- Choice of parameters :



- μ_j govern the locations of the basis functions. Can be an arbitrary set of points within the range of data
- σ governs the spatial scale. Could be chosen from the data set, e.g. average variance



Example with Gaussian Radial Basis Functions

$$\phi_j(x) = \exp\left(-\frac{(x - \mu_j)^2}{2s^2}\right)$$

basis function for $s = 0.1$ with
the μ_j on a grid with spacing s



w_j s for
middle
model:

6856.5
-3544.1
-2473.7
-2859.8
-2637.7
-2861.5
-2468.0
-3558.4

Sigmoidal Basis Functions

- Basis function of the form:


$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

where

$$\sigma(\alpha) = \frac{1}{1 + \exp(-\alpha)}$$

- Equivalently we can use \tanh :

$$\tanh(\alpha) = 2\sigma(\alpha) - 1$$



Logistic Sigmoid for
different μ_j

Other Basis Functions

- **Fourier**
 - Expansion in sinusoidal functions
 - Infinite spatial extent
- **Wavelets**
 - Localized functions in space and time
 - They have many applications in signal processing

Relationship between Maximum Likelihood and Least Squares

Goal: To show that minimizing the sum-of-squared error function is equivalent to maximizing the likelihood function under the Gaussian noise model


- Assume scalar target variable t given by:

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon$$

- $y(\mathbf{x}, \mathbf{w})$ deterministic function
- ϵ zero mean Gaussian noise

- Distribution of t is univariate normal:

$$p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1})$$



Likelihood Function

- **Dataset:**

- input values $\mathbf{X} = \{x_1, \dots, x_N\}$
- target values $\mathbf{t} = \{t_1, \dots, t_N\}$


- **Likelihood:**

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

It is the probability of observing the target data assuming they are independent

- **Log-Likelihood:** It is computationally more convenient


$$\ln p(\mathbf{t}|\mathbf{w}, \beta) = \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$



Likelihood Function

For the univariate Gaussian we have:

$$\begin{aligned}\ln p(\mathbf{t}|\mathbf{w}, \beta) &= \\ &= \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w})\end{aligned}$$



where the sum-of-squares error function is defined by:

$$E_D(\mathbf{w}) = \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2$$

Therefore, maximizing the likelihood is equivalent to minimizing the error function $E_D(\mathbf{w})$!

Maximum Likelihood Solution

We take the gradient of the log-likelihood (or of the $E_D(\mathbf{w})$):

$$\nabla \ln p(\mathbf{t}|\mathbf{w}, \beta) = \nabla E_D(\mathbf{w}) = \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\} \phi(\mathbf{x}_n)^T$$

and set it equal to zero:

$$0 = \sum_{n=1}^N t_n \phi(\mathbf{x}_n)^T - \mathbf{w}^T \left(\sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T \right)$$

and solving for \mathbf{w} we obtain:

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

That is, we can get a closed form solution(!), known as the *normal equations* for the least squares problem

The Density Matrix Φ

- Φ is an $N \times M$ matrix:

$$\Phi = \underbrace{\begin{bmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \dots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \dots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & & & \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \dots & \phi_{M-1}(\mathbf{x}_N) \end{bmatrix}}_{M\text{-basis functions}}$$

whose elements $\phi_j(\mathbf{x}_n)$ are M -basis functions, e.g. j^{th} order polynomial, or Gaussians centered on M -data points

- *Moore-Penrose pseudo-inverse* of matrix Φ

$$\Phi^\dagger \equiv (\Phi^T \Phi)^{-1} \Phi^T$$

Maximum Likelihood for precision β

Again, starting from the log-likelihood function:

$$\ln p(\mathbf{t}|\mathbf{w}, \beta) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w})$$

and maximizing with respect to the noise precision parameter β we get:

$$\frac{1}{\beta_{ML}} = \frac{1}{N} \sum_{n=1}^N \left\{ t_n - \mathbf{w}_{ML}^T \phi(\mathbf{x}_n) \right\}^2$$

Thus, the inverse of the noise precision is given by the residual variance of the target values around the regression function.


Difficulty of Direct Solution

- Direct solution of the normal equations:

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- Direct numerical calculation of the closed-form solution is not always feasible
 - when $(\Phi^T \Phi)^{-1}$ is (close) to singular (i.e. $\det(\Phi) \simeq 0$)
 - when two basis functions are colinear, the resulting parameters values can have large magnitudes
- This is not uncommon with real data sets
- It can be addressed using:
 - Singular Value Decomposition
 - Addition of regularization term can ensure non-singular matrix
 - Applying alternative numerical methods (e.g. gradient-descent)

Method of Gradient Descent




- Criterion $f(x)$ is minimized by moving from current solution in the direction of negative of gradient $f'(x)$
- Steepest descent proposes a new point

$$x' = x - \eta f'(x)$$

where $\eta > 0$ is the *learning rate*, scalar set to a small constant

Gradient with multiple inputs

Direction in w_0 - w_1 plane producing steepest descent



- For multiple inputs we need partial derivatives.
- $\frac{\partial}{\partial x_i} f(\mathbf{x})$ is how f changes as only x_i increases, gradient of f is a vector of partial derivatives $\nabla_{\mathbf{x}} f(\mathbf{x})$
- Gradient descent proposes a new point:

$$\mathbf{x}' = \mathbf{x} - \eta \nabla_{\mathbf{x}} f(\mathbf{x})$$

where $\eta > 0$ is the *learning rate*, scalar set to a small constant

Stochastic Gradient Descent for Regression

- Error function $E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2$ sums over data. Denoting $E_D(\mathbf{w}) = \sum_n E_n$, we update \mathbf{w} using

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n$$

where τ is the iteration number and η the learning rate parameter
item Replacing for the $\nabla E_n = -\{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\} \phi(\mathbf{x}_n)^T$ we get:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta(t_n - \mathbf{w}^{(\tau)T} \phi_n) \phi_n^T$$



where

- \mathbf{w} is initialized to some starting vector \mathbf{w}^0
- η chosen with care to ensure convergence
- Known as **Least Mean Squares Algorithm**

Sequential (on-line) Learning

- Maximum likelihood solution is : $\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$
- It is a batch technique
 - Processing entire training set in one go
 - Computationally expensive for large data sets,
due to huge $N \times M$ design matrix Φ
- **Solution:** use a sequential algorithm where samples are presented one at a time (or a minibatch at a time) – called *stochastic gradient descent - SGD*
- **Computational bottleneck:** is a recurring problem in ML:
 - large training sets are necessary for good generalization
 - but large training sets are also computationally expensive
- SGD is an extension of gradient descent offering a solution
 - moreover it is a method of generalization beyond the training set

Regularized Least Squares



- As model complexity increases, e.g. degree of polynomial or number of basis function then we are likely to have overfitting
- We can control overfitting by adding a regularization term to the error function

$$E(\mathbf{w}) = E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

where λ is the *regularization coefficient*

Simplest regularization form (*weight-decay*)

- The simplest regularization form is the sum-of-squares of the weight vector elements:

$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

then, the total error function becomes:

$$\frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- The error function remains a *quadratic function* of \mathbf{w} , so its exact minimizer can be found in analytical form:


$$\mathbf{w} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

which is a simple extension of the least squared solution


$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

A more general regularizer


$$\sqrt{w_1} + \sqrt{w_2} = c$$




$$w_1 + w_2 = c$$



$$w_1^2 + w_2^2 = c$$



$$w_1^4 + w_2^4 = c$$



- In a more general form the regularized error becomes:

$$\frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$

- $q = 2$ corresponds to the *quadratic equalizer*
- $q = 1$ is known as *Lasso*

Geometric interpretation of regularizer

- **Unregularized case**

We are trying to find \mathbf{w} that minimizes:


$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2$$



- **Regularized case**


We choose that value of \mathbf{w} subject to the constraint

$$\sum_{j=1}^M |w_j|^q \leq \eta$$




Sparsity with Lasso constraint

Quadratic solution where w_1^* and w_2^* are nonzero



Minimization with Lasso Regularizer. A sparse solution with $w_1^* = 0$



- With $q = 1$ and λ is sufficiently large, some of the coefficients w_j are driven to zero.
- This leads to a sparse model where the corresponding basis functions play no role.

Regularization: Summary

- Regularization allows complex models to be trained on limited size datasets without severe overfitting
- The problem of determining the optimal model complexity is shifted to determining a suitable λ value

Multiple Outputs: $\mathbf{t} = (t_1, \dots, t_K)$, $K > 1$

- It can be treated as multiple independent regression problems, introducing a different set of basis functions for each component of \mathbf{t}
- A more common approach is to introduce the same set of basis functions to model the target vector \mathbf{t} components:

$$\mathbf{y}(\mathbf{x}, \mathbf{w}) = \mathbf{W}^T \boldsymbol{\phi}(\mathbf{x})$$

where \mathbf{y} is a K -dimensional column vector; \mathbf{W} is an $M \times K$ matrix of parameters; $\boldsymbol{\phi}(\mathbf{x})$ is an M -dimensional column vector with elements $\phi_j(\mathbf{x})$, and $\phi_0(\mathbf{x}) = 1$

Solution for Multiple Outputs

- We assume the target vector \mathbf{t} has a conditional distribution of the form:

$$p(\mathbf{t}|\mathbf{x}, \mathbf{W}, \beta) = \mathcal{N}(\mathbf{t}|\mathbf{W}^T \phi(\mathbf{x}), \beta^{-1} \mathbf{I})$$

- Log-Likelihood

$$\begin{aligned} \ln p(\mathbf{T}|\mathbf{X}, \mathbf{X}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(\mathbf{t}_n|\mathbf{W}^T \phi(\mathbf{x}_n), \beta^{-1} \mathbf{I}) \\ &= \frac{NK}{2} \ln \left(\frac{\beta}{2\pi} \right) - \frac{\beta}{2} \sum_{n=1}^N \|\mathbf{t}_n - \mathbf{W}^T \phi(\mathbf{x}_n)\|^2 \end{aligned}$$

where

- \mathbf{T} is $N \times K$ matrix combining the $\mathbf{t}_1, \dots, \mathbf{t}_N$ observation set
- \mathbf{X} combines the input vectors $\mathbf{x}_1, \dots, \mathbf{x}_N$

Solution for Multiple Outputs

- The solution that maximized the log-likelihood is:

$$\mathbf{W}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{T}$$

- For each target variable t_k , we have:

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}_k = \Phi^\dagger \mathbf{t}_k$$

Conclusion:

The solution decouples between the target variables, thus we need only to compute a single pseudo-inverse matrix Φ^\dagger , shared by all of the vectors \mathbf{w}_k